

# Making Hadoop TAR Aware

Joydip Datta  
July 15, 2013

---

## Executive Summary

TAR is a widely used format for storing and distributing large collections of files such as backup images, large datasets *etc.* This documents details the use of FileSystem abstraction to access a TAR file from Hadoop without extracting it.

## Motivation

[TAR](#) is a very popular tool to combine many small files and directories into one single large file. Together with compression methods such as gunzip, TAR provides a nice way to archive, store and distribute large numbers of files together. TAR is also very popular format for storing backup images, distributing large datasets *etc.* Many of those files could be used as an input to analytic jobs.

Apache Hadoop, as of now, is not TAR aware. That is, it can not directly read a file *inside* a TAR. Neither it can run map-reduce on those files. To run analytic jobs on a TAR, one needs to first copy it to local disk, un-TAR it, then copy back to Hadoop file system. Or convert it to sequence file/other Hadoop aware format using custom (java) program. This procedure is time consuming and the user ends up having two copies of data.

By making Hadoop TAR aware, Hadoop can directly read files *inside* a TAR and run analytic jobs on those file. This way, no conversion/extraction is required.

## High Level Design

We use the Hadoop [FileSystem](#) abstraction to expose the TAR file. We implement a TarFileSystem class that extends FileSystem. All storage access in Hadoop happens through an well defined interface called FileSystem. By providing a FileSystem for TAR, virtually all Hadoop applications (including [FS-Shell](#) and Mapreduce) becomes Hadoop Aware. The implementation is mostly modeled after Hadoop Archives and employs similar access mechanisms.

## Using TAR File System

Hadoop can access a TAR archive using TAR URI SCHEMA (URI starting with `tar://`). The following examples shows this:

## Following is a TAR inside Hadoop file System

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop fs -ls /tardemo/archive.tar ↵
Found 1 items
-rw-r--r-- 1 jd supergroup 1751040 2013-07-15 20:30 /tardemo/archive.tar
```

## To access files inside this tar, simply prepone this with tar:// to make it a TAR File System URI

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop fs -ls tar:///tardemo/archive.tar ↵
13/07/15 20:33:04 INFO tar.TarFileSystem: *** Using Tar file system ***
Found 3 items
-rw-rw-r-- 1 jd jd 502760 2013-07-15 20:27 /tardemo/archive.tar+data/file2.txt
-rw-rw-r-- 1 jd jd 594933 2013-07-15 20:26 /tardemo/archive.tar+data/file1.txt
-rw-rw-r-- 1 jd jd 641720 2013-07-15 20:27 /tardemo/archive.tar+data/file3.txt
```

## To access a file inside a TAR archive, append the name of the file after the TAR URI using a '+' sign

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop fs -cat
tar:///tardemo/archive.tar+data/file1.txt ↵
13/07/15 20:38:35 INFO tar.TarFileSystem: *** Using Tar file system ***
This is the file content.
[...]
```

In TAR File System, the TAR archive is modeled like a directory and all the files inside a TAR are modeled like files within a directory. One can run mapreduce jobs on files within a TAR archive just like they do it on normal files.

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop jar hadoop*examples*.jar wordcount ↵
tar:///tardemo/archive.tar wc_out
13/07/15 20:43:05 INFO tar.TarFileSystem: *** Using Tar file system ***
13/07/15 20:43:05 INFO input.FileInputFormat: Total input paths to process : 3
13/07/15 20:43:05 INFO mapred.JobClient: Running job: job_201307151954_0001
13/07/15 20:43:06 INFO mapred.JobClient: map 0% reduce 0%
[...]
```

## To access a file inside a TAR archive, append the name of the file after the TAR URI using a '+' sign

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop fs -cat
tar:///hdfs-localhost:54310/tardemo/archive.tar+data/file1.txt ↵
13/07/15 20:38:35 INFO tar.TarFileSystem: *** Using Tar file system ***
This is the file content.
[...]
```

To access files in a different Hadoop cluster, you may mention the underlying FS schema (the FS that hosts the TAR file), NameNode address and port number in the URI

```
[jd@morpheus hadoop-1.0.3]$ bin/hadoop fs -ls
tar://hdfs-localhost:54310/tardemo/archive.tar
13/07/15 20:35:10 INFO tar.TarFileSystem: *** Using Tar file system ***
Found 3 items
-rw-rw-r--  1 jd jd      502760 2013-07-15 20:27 /tardemo/archive.tar+data/file2.txt
-rw-rw-r--  1 jd jd      594933 2013-07-15 20:26 /tardemo/archive.tar+data/file1.txt
-rw-rw-r--  1 jd jd      641720 2013-07-15 20:27 /tardemo/archive.tar+data/file3.txt
```

## Distribution and Configuration

TAR File System binary for Hadoop is distributed as a JAR library (TarFileSystem.jar). This JAR contains the main `TarFileSystem` class and other supporting classes. The user needs to copy this JAR to the `HADOOP_HOME/lib` directory or add this to `HADOOP_CLASSPATH` environment variable. You also need to expose `tar://` uri schema to Hadoop by adding the following property in `HADOOP_CONF_DIR/core-site.xml`

```
<property>
  <name>fs.tar.impl</name>
  <value>org.apache.hadoop.fs.tar.TarFileSystem</value>
</property>
```

## Conclusion

TAR is a widely used format for storing and distributing large collections of files such as backup images, large datasets *etc.* Using TarFileSystem plugin for Hadoop, one can list, read and run analytic jobs on TAR images without the need of extraction or conversion.

Hadoop also performs badly for large number of small files. To make get around this and to get the performance gain from sequential reads, one needs to create Hadoop Archives or Sequence files. But those formats are not understood outside Hadoop. As a second benefit of TAR File System, we can create one split for multiple consecutive small files to speed up mapreduce jobs.